

# Algorithms Lab

October 3, 2012

## Outline

- 1 Hints
- 2 Templates
- 3 Memory layout
- 4 Performance

Navigation icons

Algorithms Lab

## Outline

- 1 Hints
- 2 Templates
- 3 Memory layout
- 4 Performance

Navigation icons

Algorithms Lab

## Exercise 1 - Longest Path

- Modify DFS/BFS – for a given vertex  $v$ , find the longest path which starts in  $v$
- Iterate over all  $v - \mathcal{O}(n^2)$  in total ... too slow
- Modify DFS further
- Challenge – implement without a recursion

Navigation icons

Algorithms Lab

## Exercise 1 - Longest Path

Know standard techniques

- backtrack
- greedy
- divide & conquer
- dynamic programming ...

Navigation icons

Algorithms Lab

- Modify DFS/BFS – for a given vertex  $v$ , find the longest path which starts in  $v$
- Iterate over all  $v - \mathcal{O}(n^2)$  in total ... too slow
- Modify DFS further
- Challenge – implement without a recursion

Navigation icons

Algorithms Lab

## Exercise 1 - Longest Path

- Modify DFS/BFS – for a given vertex  $v$ , find the longest path which starts in  $v$
- Iterate over all  $v - \mathcal{O}(n^2)$  in total ... too slow
- Modify DFS further
- Challenge – implement without a recursion

Algorithms Lab

## Exercise 2 - The Beasts

- 50 points – backtrack
  - keep track of occupied columns/diagonals
- 100 points – backtrack with randomization, ...
- Try not to look on wikipedia for a solution!

Algorithms Lab

## Exercise 2 - The Beasts

- 50 points – backtrack
  - keep track of occupied columns/diagonals
- 100 points – backtrack with randomization, ...
- Try not to look on wikipedia for a solution!

Algorithms Lab

## Exercise 2 - The Beasts

- 50 points – backtrack
  - keep track of occupied columns/diagonals
- 100 points – backtrack with randomization, ...
- Try not to look on wikipedia for a solution!

Algorithms Lab

### Exercise 3 - Burning Coins

- 2 coins - easy
- 3 coins - still easy
- 4+ coins - draw a tree of all possible game scenarios
- Straightforward recursive implementation – 30 points
- Observation - (left, right, left) and (left, left, right) give the same resulting sequence
- Challenge - implement without a recursion

Algorithms Lab

### Exercise 3 - Burning Coins

- 2 coins - easy
- 3 coins - still easy
- 4+ coins - draw a tree of all possible game scenarios
- Straightforward recursive implementation – 30 points
- Observation - (left, right, left) and (left, left, right) give the same resulting sequence
- Challenge - implement without a recursion

Algorithms Lab

### Exercise 3 - Burning Coins

- 2 coins - easy
- 3 coins - still easy
- 4+ coins - draw a tree of all possible game scenarios
- Straightforward recursive implementation – 30 points
- Observation - (left, right, left) and (left, left, right) give the same resulting sequence
- Challenge - implement without a recursion

Navigation icons

Algorithms Lab

### Exercise 3 - Burning Coins

- 2 coins - easy
- 3 coins - still easy
- 4+ coins - draw a tree of all possible game scenarios
- Straightforward recursive implementation – 30 points
- Observation - (left, right, left) and (left, left, right) give the same resulting sequence
- Challenge - implement without a recursion

Navigation icons

Algorithms Lab

### Exercise 3 - Burning Coins

Exam question!

- $\approx 110$  students
- 31 - 100 points
- 16 - 30 points

Navigation icons

Algorithms Lab

### Exercise 4 - Pinball

- In disguise

$$\max_{i < j} \left\{ \sum_{t=1}^{i-1} p_t + (j-i+1)p_j + \sum_{t=j+1}^n p_t \right\}$$

- equivalently

$$\max_{i < j} \{S_n - (S_j - S_{i-1}) + (j-i+1)p_j\}$$

- simple  $\mathcal{O}(n^2)$  solution - iterate over all  $i, j$  – 10 points

Navigation icons

Algorithms Lab

### Exercise 4 - Pinball

- In disguise

$$\max_{i < j} \left\{ \sum_{t=1}^{i-1} p_t + (j-i+1)p_j + \sum_{t=j+1}^n p_t \right\}$$

- equivalently

$$\max_{i < j} \{S_n - (S_j - S_{i-1}) + (j-i+1)p_j\}$$

- simple  $\mathcal{O}(n^2)$  solution - iterate over all  $i, j$  – 10 points

Navigation icons

Algorithms Lab

### Exercise 4 - Pinball

- In disguise

$$\max_{i < j} \left\{ \sum_{t=1}^{i-1} p_t + (j-i+1)p_j + \sum_{t=j+1}^n p_t \right\}$$

- equivalently

$$\max_{i < j} \{S_n - (S_j - S_{i-1}) + (j-i+1)p_j\}$$

- simple  $\mathcal{O}(n^2)$  solution - iterate over all  $i, j$  – 10 points

Navigation icons

Algorithms Lab

## Exercise 4 - Pinball

For each  $j$

- $f_j(i) := S_n - (S_j - S_{i-1}) + (j - i + 1)p_j$
- $m_j$  – smallest  $i$  which maximizes  $f_j$

### Observation 1

- Consider  $j_1 < j_2$
- Assume  $p_{j'} \leq p_{j_2}$  for each  $j' \in \{j_1, \dots, j_2 - 1\}$
- then  $f_{j_2}(m_{j_1}) \geq f_{j_1}(m_{j_1}) \Rightarrow f_{j_2}(m_{j_2}) \geq f_{j_1}(m_{j_1})$

Algorithms Lab

## Exercise 4 - Pinball

For each  $j$

- $f_j(i) := S_n - (S_j - S_{i-1}) + (j - i + 1)p_j$
- $m_j$  – smallest  $i$  which maximizes  $f_j$

### Observation 1

- Consider  $j_1 < j_2$
- Assume  $p_{j'} \leq p_{j_2}$  for each  $j' \in \{j_1, \dots, j_2 - 1\}$
- then  $f_{j_2}(m_{j_1}) \geq f_{j_1}(m_{j_1}) \Rightarrow f_{j_2}(m_{j_2}) \geq f_{j_1}(m_{j_1})$

Algorithms Lab

## Exercise 4 - Pinball

For each  $j$

- $f_j(i) := S_n - (S_j - S_{i-1}) + (j - i + 1)p_j$
- $m_j$  – smallest  $i$  which maximizes  $f_j$

### Observation 1

- Consider  $j_1 < j_2$
- Assume  $p_{j'} \leq p_{j_2}$  for each  $j' \in \{j_1, \dots, j_2 - 1\}$
- then  $f_{j_2}(m_{j_1}) \geq f_{j_1}(m_{j_1}) \Rightarrow f_{j_2}(m_{j_2}) \geq f_{j_1}(m_{j_1})$

Algorithms Lab

## Exercise 4 - Pinball

For each  $j$

- $f_j(i) := S_n - (S_j - S_{i-1}) + (j - i + 1)p_j$
- $m_j$  – smallest  $i$  which maximizes  $f_j$

### Observation 1

- Consider  $j_1 < j_2$
- Assume  $p_{j'} \leq p_{j_2}$  for each  $j' \in \{j_1, \dots, j_2 - 1\}$
- then  $f_{j_2}(m_{j_1}) \geq f_{j_1}(m_{j_1}) \Rightarrow f_{j_2}(m_{j_2}) \geq f_{j_1}(m_{j_1})$

Algorithms Lab

## Exercise 4 - Pinball

Instead of every  $j$  - set of candidates  $S$  for the right end of your obstacle

- go from right to left
- consider a hole only if it updates the current maximum
- for each  $j \in S$  - find  $m_j$  by iterating  $i = 1, \dots, j - 1$
- $S$  can be large ... still  $\mathcal{O}(n^2)$  – 30 points

Algorithms Lab

## Exercise 4 - Pinball

Instead of every  $j$  - set of candidates  $S$  for the right end of your obstacle

- go from right to left
- consider a hole only if it updates the current maximum
- for each  $j \in S$  - find  $m_j$  by iterating  $i = 1, \dots, j - 1$
- $S$  can be large ... still  $\mathcal{O}(n^2)$  – 30 points

Algorithms Lab

## Exercise 4 - Pinball

### Observation 2

- Consider  $j_1, j_2 \in S$  and  $j_1 < j_2$
- remember – from the definition of  $S$  we have  $p_{j_1} > p_{j_2}$
- then  $m_{j_1} \leq m_{j_2}$

### Consequence of observation 2

- Consider some  $j \in S$  and  $m_j$
- for all  $j_1 \in S, j_1 < j : m_{j_1} \leq m_j$
- for all  $j_2 \in S, j < j_2 : m_j \leq m_{j_2}$
- what if  $j$  is the median of  $S$ ?

Algorithms Lab

## Exercise 4 - Pinball

### Observation 2

- Consider  $j_1, j_2 \in S$  and  $j_1 < j_2$
- remember – from the definition of  $S$  we have  $p_{j_1} > p_{j_2}$
- then  $m_{j_1} \leq m_{j_2}$

### Consequence of observation 2

- Consider some  $j \in S$  and  $m_j$
- for all  $j_1 \in S, j_1 < j : m_{j_1} \leq m_j$
- for all  $j_2 \in S, j < j_2 : m_j \leq m_{j_2}$
- what if  $j$  is the median of  $S$ ?

Algorithms Lab

## Exercise 4 - Pinball

### Observation 2

- Consider  $j_1, j_2 \in S$  and  $j_1 < j_2$
- remember – from the definition of  $S$  we have  $p_{j_1} > p_{j_2}$
- then  $m_{j_1} \leq m_{j_2}$

### Consequence of observation 2

- Consider some  $j \in S$  and  $m_j$
- for all  $j_1 \in S, j_1 < j : m_{j_1} \leq m_j$
- for all  $j_2 \in S, j < j_2 : m_j \leq m_{j_2}$
- what if  $j$  is the median of  $S$ ?

Algorithms Lab

## Exercise 4 - Pinball

### Observation 2

- Consider  $j_1, j_2 \in S$  and  $j_1 < j_2$
- remember – from the definition of  $S$  we have  $p_{j_1} > p_{j_2}$
- then  $m_{j_1} \leq m_{j_2}$

### Consequence of observation 2

- Consider some  $j \in S$  and  $m_j$
- for all  $j_1 \in S, j_1 < j : m_{j_1} \leq m_j$
- for all  $j_2 \in S, j < j_2 : m_j \leq m_{j_2}$
- what if  $j$  is the median of  $S$ ?

Algorithms Lab

## Exercise 4 - Pinball

### Observation 2

- Consider  $j_1, j_2 \in S$  and  $j_1 < j_2$
- remember – from the definition of  $S$  we have  $p_{j_1} > p_{j_2}$
- then  $m_{j_1} \leq m_{j_2}$

### Consequence of observation 2

- Consider some  $j \in S$  and  $m_j$
- for all  $j_1 \in S, j_1 < j : m_{j_1} \leq m_j$
- for all  $j_2 \in S, j < j_2 : m_j \leq m_{j_2}$
- what if  $j$  is the median of  $S$ ?

Algorithms Lab

## Testing hints

- Write a slow solution – easy to check if correct
- Handcraft testcases
  - whatever comes to your mind
  - trivial cases – 5, 4, 3, 2, 1 for Pinball
- Generate testcases
  - random

Algorithms Lab

## Testing hints

- Write a slow solution – easy to check if correct
- Handcraft testcases
  - whatever comes to your mind
  - trivial cases – 5, 4, 3, 2, 1 for Pinball
- Generate testcases
  - random

## Testing hints

- Write a slow solution – easy to check if correct
- Handcraft testcases
  - whatever comes to your mind
  - trivial cases – 5, 4, 3, 2, 1 for Pinball
- Generate testcases
  - random

## Testing hints

### Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Testing hints

### Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Testing hints

### Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Testing hints

### Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Testing hints

Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Testing hints

Random testcases

- might miss border cases – correctness
- might miss worse case – runtime

### Example

- Consider second solution for the Pinball problem
- Runtime –  $\mathcal{O}(|S| \cdot n)$
- Generate random testcase – assume all points are distinct

$$\Pr[\text{update of maximum at position } i] = \frac{1}{n - i + 1}$$

$$\mathbb{E}[\# \text{ of updates}] = \sum_{i=1}^n 1/i \approx \log n$$

- $|S| \approx \log n \rightarrow \mathcal{O}(n \log n)$  running time in expectation

## Outline

- 1 Hints
- 2 Templates
- 3 Memory layout
- 4 Performance

## Templates

### Example

```
std::vector<int>
```

## Templates – goal

Templates solve this problem:

```
1 void swap_int (int& a, int& b) {
2     int t = b;
3     b = a;
4     a = t;
5 }
```

## Templates – goal

Templates solve this problem:

```
1 void swap_int (int& a, int& b) {
2     int t = b;
3     b = a;
4     a = t;
5 }
6 void swap_float (float& a, float& b) {
7     float t = b;
8     b = a;
9     a = t;
10 }
```

## Function templates

How about if we could say

```
1  /* for any type T, define a function */
2  void swap (T& a, T& b) {
3      T t = b;
4      b = a;
5      a = t;
6  }
```

It is that easy!

```
1  template<class T>
2  void swap (T& a, T& b) {
3      T t = b;
4      b = a;
5      a = t;
6  }
```

Instead of class, you can also use typename as a synonym.

Navigation icons

Algorithms Lab

## Function templates

How about if we could say

```
1  /* for any type T, define a function */
2  void swap (T& a, T& b) {
3      T t = b;
4      b = a;
5      a = t;
6  }
```

It is that easy!

```
1  template<class T>
2  void swap (T& a, T& b) {
3      T t = b;
4      b = a;
5      a = t;
6  }
```

Instead of class, you can also use typename as a synonym.

Navigation icons

Algorithms Lab

## Class templates

```
1  template<class T, class U>
2  class pair {
3  public:
4      T first;
5      U second;
6  };
```

```
7  pair<int, float> p;
8  p.first = 1;
9  p.second = 2;
```

The same goes for struct.

Navigation icons

Algorithms Lab

## Class templates

```
1  template<class T, class U>
2  class pair {
3  public:
4      T first;
5      U second;
6  };
```

```
7  pair<int, float> p;
8  p.first = 1;
9  p.second = 2;
```

The same goes for struct.

Navigation icons

Algorithms Lab

## Special-casing types

Templates can special-case for certain types (or even values).  
Consider a simple vector:

```
1  template<class T>
2  class vector {
3      long _size;
4      long _capacity;
5      T *data;
6  public:
7      T operator[](int i) const {
8          return data[i];
9      }
10     /* rest of logic left as an exercise */
11 };
```

But if  $T$  is `bool`, shouldn't we implement a packed bitfield?

Navigation icons

Algorithms Lab

## Special-casing types

Templates can special-case for certain types (or even values).  
Consider a simple vector:

```
1  template<class T>
2  class vector {
3      T *data;
4  };
```

But if  $T$  is `bool`, shouldn't we implement a packed bitfield?

Navigation icons

Algorithms Lab



## Special-cased bitfield vector

```
1 template<>
2 class vector<bool> {
3     long _size;
4     long _capacity;
5     char *data;
6 public:
7     bool operator[](int i) const {
8         return (data[i/8] >> i%8) & 1;
9     }
10    /* rest of logic left as an exercise */
11 };
```

## Outline

- 1 Hints
- 2 Templates
- 3 Memory layout
- 4 Performance

## Memory layout and variables

```
1 int a;
2 static int b;
3 extern int c;
4 std::vector<int> u(1);
5
6 int f(int x) {
7     int i;
8     static int j;
9     std::vector<int> v;
10    for (int k = 0; k < 100; k++) {
11        std::vector<int> w(1);
12    }
13 }
```

## Outline

- 1 Hints
- 2 Templates
- 3 Memory layout
- 4 Performance

## Performance – measures

**Real time** or **wall time**: Time elapsed on a clock until the task is finished

**User time** Time the kernel allocated for your program to run

**System time** Time the kernel did work on your program's behalf

### Note

Measurement is influenced by many things and variance usually about 20%.

## Performance – measures

**Real time** or **wall time**: Time elapsed on a clock until the task is finished

**User time** Time the kernel allocated for your program to run

**System time** Time the kernel did work on your program's behalf

### Note

Measurement is influenced by many things and variance usually about 20%.

## User/system time example

Consider

```
1 for (int i = 0; i < 1000000; i++)
2     std::cout << i << "\n";
```

vs.

```
1 for (int i = 0; i < 1000000; i++)
2     std::cout << i << std::endl;
```

	real [ms]	user [ms]	sys [ms]
"\n"	99	97	2
endl	241	123	117

## User/system time example

Consider

```
1 for (int i = 0; i < 1000000; i++)
2     std::cout << i << "\n";
```

vs.

```
1 for (int i = 0; i < 1000000; i++)
2     std::cout << i << std::endl;
```

	real [ms]	user [ms]	sys [ms]
"\n"	99	97	2
endl	241	123	117

## Measurement on judge

### Speed

The judge measures **only** something vaguely resembling **user time**.

- For the purposes of this tutorial we call this unit *ticks* (t).
- The "actual" time measurements were done on my laptop.
  - Clock speed boosted 3.3 GHz
  - Length of one clock cycle thus 0.3 ns

## Measurement on judge

### Speed

The judge measures only something vaguely resembling user time.

- For the purposes of this tutorial we call this unit *ticks* (t).
- The "actual" time measurements were done on my laptop.
  - Clock speed boosted 3.3 GHz
  - Length of one clock cycle thus 0.3 ns

## Measurement on judge

### Speed

The judge measures only something vaguely resembling user time.

- For the purposes of this tutorial we call this unit *ticks* (t).
- The "actual" time measurements were done on my laptop.
  - Clock speed boosted 3.3 GHz
  - Length of one clock cycle thus 0.3 ns

## Measurement on judge

### Speed

The judge measures only something vaguely resembling user time.

- For the purposes of this tutorial we call this unit *ticks* (t).
- The "actual" time measurements were done on my laptop.
  - Clock speed boosted 3.3 GHz
  - Length of one clock cycle thus 0.3 ns

## Performance: simple operations

X-ing up a 1000-element array takes (per element)...

op	ns	nt
add	0.17	0.26
multiply	0.56	0.52
divide	6.0	1.2

Remember that one clock cycle is 510.3ns!

## Performance: input/output

Doing I/O with `int` typed 1–3 digit numbers takes...

op	ns	nt
cin sync	240	297
cin nosync	79	100
scanf	95	124
cout sync	84	118
cout nosync	90	107
cout sync endl	130	142
cout nosync endl	110	120
printf	81	106

## Performance: function calls

Calling a simple function...

op	ns	nt
directly	1.9	2.1
indirectly	2.2	2.3
inlined	0.23	0.28

## Performance: memory allocations

Allocating a few bytes of memory on the heap...

op	ns	nt
<code>new</code>	31	41
<code>new, delete</code>	35	52
<code>malloc</code>	27	38
<code>malloc, free</code>	32	48

## Performance: vector

Storing elements in a vector of type...

type	ns	nt
<code>int</code>	0.23	0.20
<code>char</code>	0.28	0.35
<code>bool</code>	0.89	1.65

## Performance: set

In a set, doing...

op	ns	nt
insert	180	130
insert with many dups	90	31
insert, remove	280	65

Recall that `new` is about 30 ns.

Performance: vector

Sum up the elements in a... (per element)

type	ns	nt
vector<int>	0.6	0.85
set<int>	5.9	4.2

Algorithms Lab

Performance: vector

Sum up the elements in a... (per element)

type	ns	nt
vector<int>	0.6	0.85
set<int>	5.9	4.2

Algorithms Lab

Performance: synchronization

Locking and then unlocking a mutex...

type	ns
only one user	4.7
heavily contended	140

Algorithms Lab

Performance: summary

Algorithms Lab